

```
!pip install -q pyathena
```

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

```
import boto3
import sagemaker
import pandas as pd
from pyathena import connect
```

```
sagemaker.config INFO - Not applying SDK defaults from location:
/etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location:
/root/.config/sagemaker/config.yaml
```

## Auth with AWS

```
sess = sagemaker.Session()
bucket = sess.default_bucket()
role = sagemaker.get_execution_role()
region = boto3.Session().region_name
account_id = boto3.client("sts").get_caller_identity().get("Account")

sm = boto3.Session().client(service_name="sagemaker",
region_name=region)
```

## Dropping album\_name from dataset

```
raw_df = pd.read_csv('dataset.csv')
display(raw_df.head())
```

	Unnamed: 0	track_id	artists	
0	0	5Su0ikwiRyPMVoIQDJUgSV	Gen Hoshino	
1	1	4qPNDBW1i3p13qLCt0Ki3A	Ben Woodward	
2	2	1iJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN	
3	3	6lfxq3CG4xtTiEg7opyCyx	Kina Grannis	
4	4	5vjLSffimiIP26QG5WcN2K	Chord Overstreet	

  

	album_name	
0	Comedy	
1	Ghost (Acoustic)	
2	To Begin Again	
3	Crazy Rich Asians (Original Motion Picture Sou...	
4	Hold On	

  

	track_name	popularity	duration_ms	explicit	
0	Comedy	73	230666	False	
1	Ghost - Acoustic	55	149610	False	
2	To Begin Again	57	210826	False	

3	Can't Help Falling In Love	71	201933	False
4	Hold On	82	198853	False

	danceability	energy	...	loudness	mode	speechiness
acousticness \						
0	0.676	0.4610	...	-6.746	0	0.1430
0.0322						
1	0.420	0.1660	...	-17.235	1	0.0763
0.9240						
2	0.438	0.3590	...	-9.734	1	0.0557
0.2100						
3	0.266	0.0596	...	-18.515	1	0.0363
0.9050						
4	0.618	0.4430	...	-9.681	1	0.0526
0.4690						

	instrumentalness	liveness	valence	tempo	time_signature
track_genre					
0	0.000001	0.3580	0.715	87.917	4
acoustic					
1	0.000006	0.1010	0.267	77.489	4
acoustic					
2	0.000000	0.1170	0.120	76.332	4
acoustic					
3	0.000071	0.1320	0.143	181.740	3
acoustic					
4	0.000000	0.0829	0.167	119.949	4
acoustic					

[5 rows x 21 columns]

```
cleaned_df = raw_df.drop(columns=['album_name', 'Unnamed: 0'])
display(cleaned_df.head())
```

	track_id	artists
track_name \		
0	5Su0ikwiRyPMVoIQDJUgSV	Gen Hoshino
Comedy		
1	4qPNDBW1i3p13qLCt0Ki3A	Ben Woodward
Acoustic		
2	1iJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN
Begin Again		
3	6lfxq3CG4xtTiEg7opyCyx	Kina Grannis
In Love		
4	5vjLSffimiIP26QG5WcN2K	Chord Overstreet
Hold On		

	popularity	duration_ms	explicit	danceability	energy	key
loudness \						
0	73	230666	False	0.676	0.4610	1
-						

6.746							
1	55	149610	False	0.420	0.1660	1	-
17.235							
2	57	210826	False	0.438	0.3590	0	-
9.734							
3	71	201933	False	0.266	0.0596	0	-
18.515							
4	82	198853	False	0.618	0.4430	2	-
9.681							

	mode	speechiness	acousticness	instrumentalness	liveness
valence \					
0	0	0.1430	0.0322	0.000001	0.3580
0.715					
1	1	0.0763	0.9240	0.000006	0.1010
0.267					
2	1	0.0557	0.2100	0.000000	0.1170
0.120					
3	1	0.0363	0.9050	0.000071	0.1320
0.143					
4	1	0.0526	0.4690	0.000000	0.0829
0.167					

	tempo	time_signature	track_genre
0	87.917	4	acoustic
1	77.489	4	acoustic
2	76.332	4	acoustic
3	181.740	3	acoustic
4	119.949	4	acoustic

```
# create local version
cleaned_df.to_csv('dataset_clean.csv', index=False)
```

Convert csv to tsv and move to S3

```
s3_private_data_path = "s3://{}/w2-musicData/csv".format(bucket)
print(s3_private_data_path)

s3://sagemaker-us-east-1-106006112223/w2-musicData/csv

!aws s3 cp "dataset_clean.csv" $s3_private_data_path/

upload: ./dataset_clean.csv to
s3://sagemaker-us-east-1-106006112223/w2-musicData/csv/dataset_clean.c
sv

!aws s3 ls $s3_private_data_path/

2024-09-17 05:56:57    16931936 dataset_clean.csv
```

## Create DB in Athena for queries

```
# Set S3 staging directory -- this is a temporary directory used for
Athena queries
database_name = "w2_music_db"
table_name_tsv = 'music_ds_tsv10'
s3_staging_dir = "s3://{0}/athena/staging".format(bucket)
print(s3_staging_dir)
conn = connect(region_name=region, s3_staging_dir=s3_staging_dir)

s3://sagemaker-us-east-1-106006112223/athena/staging

statement = "CREATE DATABASE IF NOT EXISTS {}".format(database_name)
pd.read_sql(statement, conn)

statement = "SHOW DATABASES"
df_show = pd.read_sql(statement, conn)
df_show.head(5)

/tmp/ipykernel_140/3245868569.py:2: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
    pd.read_sql(statement, conn)
/tmp/ipykernel_140/3245868569.py:5: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
    df_show = pd.read_sql(statement, conn)

  database_name
0      default
1      dsoaws
2    w2_music_db
```

## Create tables in DB and schemas

```
# SQL statement to execute
statement = """CREATE EXTERNAL TABLE IF NOT EXISTS {}.{}(
    track_id string,
    artists string,
    track_name string,
    popularity int,
    duration_ms int,
    explicit string,
    danceability float,
    energy float,
    key int,
    loudness float,
    mode int,
    speechiness float,
```

```

        acousticness float,
        instrumentalness float,
        liveness float,
        valence float,
        tempo float,
        time_signature int,
        track_genre string
    ) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY
    '\\n' LOCATION '{}
TBLPROPERTIES ('compressionType'='gzip',
'skip.header.line.count'='1')""".format(
    database_name, table_name_tsv, s3_private_data_path
)

```

```
print(statement)
```

```
pd.read_sql(statement, conn)
```

```

CREATE EXTERNAL TABLE IF NOT EXISTS w2_music_db.music_ds_tsv10(
    track_id string,
    artists string,
    track_name string,
    popularity int,
    duration_ms int,
    explicit boolean,
    danceability float,
    energy float,
    key int,
    loudness float,
    mode int,
    speechiness float,
    acousticness float,
    instrumentalness float,
    liveness float,
    valence float,
    tempo float,
    time_signature int,
    track_genre string
) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\\
n' LOCATION 's3://sagemaker-us-east-1-106006112223/w2-musicData/csv'
TBLPROPERTIES ('compressionType'='gzip', 'skip.header.line.count'='1')

```

```

/tmp/ipykernel_359/1675701273.py:29: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.

```

```
pd.read_sql(statement, conn)
```

```
Empty DataFrame
```

```
Columns: []
```

```
Index: []
```

```
statement = "SHOW TABLES IN W2_MUSIC_DB"
```

```
df_show = pd.read_sql(statement, conn)
```

```
df_show.head(5)
```

```
/tmp/ipykernel_359/1294112312.py:3: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
df_show = pd.read_sql(statement, conn)
```

```
      tab_name
0  music_ds_tsv10
1   music_ds_tsv8
2   music_ds_tsv9
```

```
# first test query to get all data via athena
```

```
statement = """SELECT * FROM {}.{} LIMIT 5""".format(
    database_name, table_name_tsv
)
```

```
print(statement)
```

```
sql_df = pd.read_sql(statement, conn)
```

```
display(sql_df.head(5))
```

```
SELECT * FROM w2_music_db.music_ds_tsv10 LIMIT 5
```

```
/tmp/ipykernel_359/1057200777.py:7: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
sql_df = pd.read_sql(statement, conn)
```

	track_id	artists
track_name \		
0	5Su0ikwiRyPMVoIQDJUgSV	Gen Hoshino
Comedy		
1	4qPNDBWli3p13qLCt0Ki3A	Ben Woodward
Acoustic		Ghost -
2	liJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN
Begin Again		To
3	6lfxq3CG4xtTiEg7opyCyx	Kina Grannis
In Love		Can't Help Falling
4	5vjLSffimiIP26QG5WcN2K	Chord Overstreet
Hold On		

```
popularity duration_ms explicit danceability energy key
```

loudness	\						
0	73	230666	False	0.676	0.4610	1	-
6.746							
1	55	149610	False	0.420	0.1660	1	-
17.235							
2	57	210826	False	0.438	0.3590	0	-
9.734							
3	71	201933	False	0.266	0.0596	0	-
18.515							
4	82	198853	False	0.618	0.4430	2	-
9.681							

mode	speechiness	acousticness	instrumentalness	liveness
0	0.1430	0.0322	0.000001	0.3580
1	0.0763	0.9240	0.000006	0.1010
2	0.0557	0.2100	0.000000	0.1170
3	0.0363	0.9050	0.000071	0.1320
4	0.0526	0.4690	0.000000	0.0829

tempo	time_signature	track_genre
0 87.917	4	acoustic
1 77.489	4	acoustic
2 76.332	4	acoustic
3 181.740	3	acoustic
4 119.949	4	acoustic

```
# reading local csv file using pandas
full_df = pd.read_csv('dataset_clean.csv')
full_df = full_df.dropna()
display(full_df.head())
```

track_name	track_id	artists
0 5Su0ikwiRyPMVoIQDJUgSV Comedy		Gen Hoshino
1 4qPNDBW1i3p13qLCt0Ki3A Acoustic		Ben Woodward
2 1iJBSr7s7jYXzM8EGcbK5b Begin Again	Ingrid Michaelson;ZAYN	To
3 6lfxq3CG4xtTiEg7opyCyx In Love	Kina Grannis	Can't Help Falling
4 5vjLSffimiIP26QG5WcN2K Hold On	Chord Overstreet	

popularity	duration_ms	explicit	danceability	energy	key
loudness \					
0 73	230666	False	0.676	0.4610	1 -
6.746					
1 55	149610	False	0.420	0.1660	1 -
17.235					
2 57	210826	False	0.438	0.3590	0 -
9.734					
3 71	201933	False	0.266	0.0596	0 -
18.515					
4 82	198853	False	0.618	0.4430	2 -
9.681					

  

mode	speechiness	acousticness	instrumentalness	liveness
valence \				
0 0	0.1430	0.0322	0.000001	0.3580
0.715				
1 1	0.0763	0.9240	0.000006	0.1010
0.267				
2 1	0.0557	0.2100	0.000000	0.1170
0.120				
3 1	0.0363	0.9050	0.000071	0.1320
0.143				
4 1	0.0526	0.4690	0.000000	0.0829
0.167				

  

tempo	time_signature	track_genre
0 87.917	4	acoustic
1 77.489	4	acoustic
2 76.332	4	acoustic
3 181.740	3	acoustic
4 119.949	4	acoustic

## Homework queries

1. List artist, track\_name, and popularity for songs that have a popularity greater than or equal to 99

```
statement = """SELECT artists, track_name, popularity FROM {}.{}
WHERE popularity >= 99""".format(
    database_name, table_name_tsv
)

# CAST(popularity AS INTEGER) >= 99
print(statement)

df = pd.read_sql(statement, conn)
df.head(10)
```



```
SELECT artists, track_name, popularity FROM w2_music_db.music_ds_tsv10
WHERE popularity >= 99
```

```
/tmp/ipykernel_359/2117349163.py:9: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
df = pd.read_sql(statement, conn)
```

	artists	track_name	popularity
0	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100
1	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100

```
# pandas
```

```
pd_df = full_df[full_df['popularity'] >= 99]
[['artists', 'track_name', 'popularity']]
display(pd_df.head())
```

	artists	track_name	popularity
20001	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100
51664	Bizarrap;Quevedo	Quevedo: Bzrp Music Sessions, Vol. 52	99
81051	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100

## 2. List artists with an average popularity of 92

```
statement = """SELECT artists FROM {}.{}
GROUP BY artists HAVING AVG(popularity) = 92""".format(
    database_name, table_name_tsv
)
```

```
print(statement)
```

```
df = pd.read_sql(statement, conn)
df.head(10)
```

```
SELECT artists FROM w2_music_db.music_ds_tsv9
GROUP BY artists HAVING AVG(popularity) = 92
```

```
/tmp/ipykernel_245/3210071803.py:14: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
df = pd.read_sql(statement, conn)
```

	artists
0	Harry Styles
1	Rema;Selena Gomez

```
# pandas
artists_avg_popularity = full_df.groupby('artists').filter(lambda x:
x['popularity'].mean() == 92)
display(artists_avg_popularity.head())
artists_avg_popularity_list =
artists_avg_popularity['artists'].unique()
print(artists_avg_popularity_list)
```

	track_id	artists \
81052	4LRPiXqCikLlN15c3yImP7	Harry Styles
81100	0WtM2NBVQNNJLh6scP13H8	Rema;Selena Gomez
81158	6UeLLqGLWMcVH1E5c4H7lY	Harry Styles
81205	4Dvkj6JhhA12EX05fT7y2e	Harry Styles

	track_name	popularity	duration_ms
explicit \			
81052	As It Was	95	167303
False			
81100	Calm Down (with Selena Gomez)	92	239317
False			
81158	Watermelon Sugar	89	174000
False			
81205	As It Was	92	167303
False			

	danceability	energy	key	loudness	mode	speechiness
acousticness \						
81052	0.520	0.731	6	-5.338	0	0.0557
0.342						
81100	0.801	0.806	11	-5.206	1	0.0381
0.382						
81158	0.548	0.816	0	-4.209	1	0.0465
0.122						
81205	0.520	0.731	6	-5.338	0	0.0557
0.342						

	instrumentalness	liveness	valence	tempo	time_signature \
81052	0.001010	0.311	0.662	173.930	4
81100	0.000669	0.114	0.802	106.999	4
81158	0.000000	0.335	0.557	95.390	4
81205	0.001010	0.311	0.662	173.930	4

	track_genre
81052	pop
81100	pop
81158	pop
81205	pop

```
['Harry Styles' 'Rema;Selena Gomez']
```

### 3. List the Top 10 most energetic genres

```
statement = """SELECT DISTINCT track_genre
FROM {}.{}
LIMIT 10;""".format(
    database_name, table_name_tsv
)
```

```
print(statement)
df = pd.read_sql(statement, conn)
display(df.head(10))
```

*# Error in schema / parsing, track\_genre is all messed up*

```
SELECT DISTINCT track_genre
FROM w2_music_db.music_ds_tsv10
LIMIT 10;
```

/tmp/ipykernel\_359/1041569638.py:8: UserWarning: pandas only supports SQLAlchemy connectable (engine/connection) or database string URI or sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please consider using SQLAlchemy.

```
df = pd.read_sql(statement, conn)
```

```
   track_genre
0    acoustic
1             4
2             3
3    163.99
4    124.157
5    193.395
6    afrobeat
7    alt-rock
8  alternative
9    ambient
```

```
statement = """SELECT track_genre, AVG(energy) AS avg_energy FROM {}.{}
GROUP BY track_genre
ORDER BY avg_energy DESC""".format(
    database_name, table_name_tsv
)
```

```
print(statement)
df = pd.read_sql(statement, conn)
df.head(10)
```

```
SELECT track_genre, AVG(energy) AS avg_energy FROM
w2_music_db.music_ds_tsv10
GROUP BY track_genre
ORDER BY avg_energy DESC
```

```
/tmp/ipykernel_359/4047995415.py:14: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
df = pd.read_sql(statement, conn)
```

	track_genre	avg_energy
0	0.797	1174026.0
1	0.556	691306.0
2	0.492	542000.0
3	0.45	538160.0
4	0.347	526706.0
5	0.0761	502786.0
6	0.0903	449813.0
7	0.035	440310.0
8	0.483	371160.0
9	0.147	355693.0

```
# pandas
```

```
top_energetic_genres = full_df.groupby('track_genre')
['energy'].mean().sort_values(ascending=False).head(10)
display(top_energetic_genres)
```

track_genre	
death-metal	0.931470
grindcore	0.924201
metalcore	0.914485
happy	0.910971
hardstyle	0.901246
drum-and-bass	0.876635
black-metal	0.874897
heavy-metal	0.874003
party	0.871237
j-idol	0.868677

Name: energy, dtype: float64

4. How many tracks is Bad Bunny on?

```
# SELECT COUNT(*) AS track_count
# FROM w2_music_db.tracks
# WHERE artists LIKE '%Bad Bunny%';
```

```
statement = """SELECT COUNT(*) AS track_count FROM {}.{}
WHERE artists LIKE '%Bad Bunny%'""".format(
    database_name, table_name_tsv
)
```

```
print(statement)
df = pd.read_sql(statement, conn)
print(df)
```

```
/tmp/ipykernel_359/768749324.py:11: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
df = pd.read_sql(statement, conn)
```

```
SELECT COUNT(*) AS track_count FROM w2_music_db.music_ds_tsv10
WHERE artists LIKE '%Bad Bunny%'
track_count
0          416
```

```
bad_bunny_tracks_count = full_df[full_df['artists'].str.contains('Bad
Bunny')].shape[0]
print(bad_bunny_tracks_count)
```

```
416
```

5. Show the top 10 genres in terms of popularity sorted by their most popular track

```
# SELECT track_genre, MAX(popularity) AS max_popularity
# FROM w2_music_db.tracks
# GROUP BY track_genre
# ORDER BY max_popularity DESC
# LIMIT 10;
```

```
statement = """SELECT track_genre, MAX(popularity) AS max_popularity
FROM {}.{}
GROUP BY track_genre
ORDER BY max_popularity DESC
LIMIT 10""".format(
    database_name, table_name_tsv
)
```

```
print(statement)
df = pd.read_sql(statement, conn)
df.head(10)
```

```
# noticed slight difference in return... hip hop genre got ereased?
```

```
SELECT track_genre, MAX(popularity) AS max_popularity FROM
w2_music_db.music_ds_tsv10
GROUP BY track_genre
ORDER BY max_popularity DESC
LIMIT 10
```

```
/tmp/ipykernel_359/3944904740.py:15: UserWarning: pandas only supports
SQLAlchemy connectable (engine/connection) or database string URI or
sqlite3 DBAPI2 connection. Other DBAPI2 objects are not tested. Please
consider using SQLAlchemy.
```

```
df = pd.read_sql(statement, conn)
```

	track_genre	max_popularity
0	dance	100
1	pop	100
2	latin	98
3	reggaeton	98
4	latino	98
5	edm	98
6	reggae	98
7	piano	96
8	rock	96
9	chill	93

```
# pandas
```

```
top_genres_by_popularity = full_df.groupby('track_genre')
['popularity'].max().sort_values(ascending=False).head(10)
print(top_genres_by_popularity)
```

```
track_genre
dance      100
pop        100
hip-hop     99
latin       98
edm         98
latino      98
reggaeton   98
reggae      98
rock        96
piano       96
Name: popularity, dtype: int64
```

```
%%html
```

<p><b>Shutting down your kernel for this notebook to release resources.</b></p>

<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

```
<script>
```

```
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

```
%%javascript

try {
    Jupyter.notebook.save_checkpoint();
    Jupyter.notebook.session.delete();
}
catch(err) {
    // NoOp
}
```